

# Jan 23, 20 Dask Basic

February 19, 2020

- Name: Jikhan Jeong
- Ref: <https://docs.dask.org/en/latest/install.html> (conda install dask)
- Ref; <https://datascienceschool.net/view-notebook/2282b75b2a63448087b77269885c27cb/> (Korean)
- 

**0.1 Ref: <https://towardsdatascience.com/trying-out-dask-dataframes-in-python-for-fast-data-analysis-in-parallel-aa960c18a915> (English)**

```
[2]: import pandas as pd
import numpy as np
import os
print(os.getcwd())
```

/data/cahnrs/jikhan.jeong/amazon

---

#  
Part  
1:  
Ba-  
sic

---

\*  
dask  
dataframe.compute()  
->  
to  
pan-  
das  
dataframe  
\*

Ref;  
<https://datascienceschool.net/view-notebook/2282b75b2a63448087b77269885c27cb/>  
(Korean)

---

```
[3]: import dask.dataframe as dd
from dask.diagnostics import ProgressBar
pbar = ProgressBar()
pbar.register()
```

```
[8]: %%writefile data1.csv
time,temperature,humidity
0,22,58
1,21,57
2,25,57
3,26,55
4,22,53
5,23,59
```

Overwriting data1.csv

```
[9]: ls data*.csv
```

data1.csv data2.csv data3.csv

```
[10]: df = dd.read_csv("data1.csv")
df
```

```
[10]: Dask DataFrame Structure:
           time temperature humidity
npartitions=1
           int64      int64      int64
           ...      ...      ...
Dask Name: from-delayed, 3 tasks
```

```
[11]: type(df)
```

```
[11]: dask.dataframe.core.DataFrame
```

```
[12]: df.head(1)
```

```
[#####] | 100% Completed | 0.1s
[#####] | 100% Completed | 0.1s
```

```
[12]:   time  temperature  humidity
0     0             22         58
```

```
[13]: df.temperature.mean() # task so no results
```

```
[13]: dd.Scalar<series-..., dtype=float64>
```

```
[14]: df.temperature.mean().compute()
```

```
##### | 100% Completed | 0.1s
##### | 100% Completed | 0.1s
```

```
[14]: 23.166666666666668
```

```
[15]: df.temperature.compute()
```

```
##### | 100% Completed | 0.1s
##### | 100% Completed | 0.1s
```

```
[15]: 0    22
      1    21
      2    25
      3    26
      4    22
      5    23
      Name: temperature, dtype: int64
```

```
[16]: (df.temperature*10+5).compute()
```

```
##### | 100% Completed | 0.1s
##### | 100% Completed | 0.1s
```

```
[16]: 0    225
      1    215
      2    255
      3    265
      4    225
      5    235
      Name: temperature, dtype: int64
```

```
[12]: df = df.assign(temperature = df.temperature*10)
      df.head(1)
```

```
[12]:   time  temperature  humidity
      0      0         220       58
```

```
[13]: df = df.assign(title = df.temperature.astype(str) + "jikhan is handsome")
      df.head()
```

```
[13]:   time  temperature  humidity          title
      0      0         220       58  220jikhan is handsome
      1      1         210       57  210jikhan is handsome
      2      2         250       57  250jikhan is handsome
      3      3         260       55  260jikhan is handsome
      4      4         220       53  220jikhan is handsome
```

- read a file with while-card(\*)

- count().compute()
- describe().compute()
- 

0.2 E.g. `df = dd.read_csv('data*.csv') # read data1, data2, data3`

```
[17]: %%writefile data2.csv
time,temperature,humidity
0,22,58
1,21,57
2,25,57
3,26,55
4,22,53
5,23,59
```

Overwriting data2.csv

```
[18]: %%writefile data3.csv
time,temperature,humidity
0,22,58
1,21,57
2,25,57
3,26,55
4,22,53
5,23,59
```

Overwriting data3.csv

```
[19]: ls data*.csv
```

data1.csv data2.csv data3.csv

```
[20]: df1 = dd.read_csv('data*.csv')
df1.head(12)
```

```
[#####] | 100% Completed | 0.1s
```

```
[#####] | 100% Completed | 0.1s
```

```
/opt/apps/anaconda3/19.10.0/lib/python3.7/site-
packages/dask/dataframe/core.py:5738: UserWarning: Insufficient elements for
`head`. 12 elements requested, only 6 elements available. Try passing larger
`npartitions` to `head`.
warnings.warn(msg.format(n, len(r)))
```

```
[20]:   time  temperature  humidity
0     0           22         58
1     1           21         57
```

2	2	25	57
3	3	26	55
4	4	22	53
5	5	23	59

```
[21]: len(df1)
```

```
[#####] | 100% Completed | 0.1s
[#####] | 100% Completed | 0.1s
```

```
[21]: 18
```

```
[22]: df1.compute()
```

```
[#####] | 100% Completed | 0.1s
[#####] | 100% Completed | 0.1s
```

```
[22]:
```

	time	temperature	humidity
0	0	22	58
1	1	21	57
2	2	25	57
3	3	26	55
4	4	22	53
5	5	23	59
0	0	22	58
1	1	21	57
2	2	25	57
3	3	26	55
4	4	22	53
5	5	23	59
0	0	22	58
1	1	21	57
2	2	25	57
3	3	26	55
4	4	22	53
5	5	23	59

```
[23]: df1.count()
```

```
[23]: Dask Series Structure:
npartitions=1
humidity    int64
time        ...
dtype: int64
Dask Name: dataframe-count-agg, 13 tasks
```

```
[24]: df1.count().compute()
```

```
[#####] | 100% Completed | 0.1s
[#####] | 100% Completed | 0.1s
```

```
[24]: time          18
      temperature    18
      humidity       18
      dtype: int64
```

```
[25]: df1.temperature.describe().compute()
```

```
[#####] | 100% Completed | 0.1s
[#####] | 100% Completed | 0.1s
```

```
[25]: count      18.000000
      mean       23.166667
      std        1.823055
      min       21.000000
      25%       22.000000
      50%       22.500000
      75%       24.500000
      max       26.000000
      Name: temperature, dtype: float64
```

```
[26]: type(df1)
```

```
[26]: dask.dataframe.core.DataFrame
```

```
[27]: df1_pd = df1.compute()
      type(df1_pd)
```

```
[#####] | 100% Completed | 0.1s
[#####] | 100% Completed | 0.1s
```

```
[27]: pandas.core.frame.DataFrame
```

- Dataset: <https://catalogdata.gov.uk/dataset/tucrimes2001stospresents398a4> (1.3GB)
- (Window) \* using **urllib.reques** instead of wget in Linux
- (Linux) `!wget -O crime.csv https://data.cityofchicago.org/api/views/ijzp-q8t2/rows.csv?accessType=DOWNLOAD` (Linux)
- 

0.3 Ref: <https://stackoverflow.com/questions/57748687/downloading-files-in-jupyter-wget-on-windows>

```
[24]: import urllib.request
```



267231	24560	JC279072	05/26/2019	06:48:00 AM	013XX W HASTINGS ST
267232	24559	JC278908	05/26/2019	02:11:00 AM	013XX W HASTINGS ST
267233	11707734	JC287730	07/01/2014	07:30:00 AM	063XX S NORMAL BLVD
267234	11707239	JC287563	11/30/2017	09:00:00 AM	022XX S KOSTNER AVE

	IUCR	Primary Type	Description \
267230	2825	OTHER OFFENSE	HARASSMENT BY TELEPHONE
267231	0110	HOMICIDE	FIRST DEGREE MURDER
267232	0110	HOMICIDE	FIRST DEGREE MURDER
267233	1153	DECEPTIVE PRACTICE	FINANCIAL IDENTITY THEFT OVER \$ 300
267234	1153	DECEPTIVE PRACTICE	FINANCIAL IDENTITY THEFT OVER \$ 300

	Location	Description	Arrest	Domestic	...	Ward	Community	Area	FBI Code \
267230		APARTMENT	false	true	...	24		29	26
267231		CHA PARKING LOT	true	false	...	25		28	01A
267232		STREET	false	false	...	25		28	01A
267233		NaN	false	false	...	20		68	11
267234		RESIDENCE	false	false	...	22		29	11

	X Coordinate	Y Coordinate	Year	Updated On	Latitude \
267230	1152126	1893208	2019	06/30/2019 03:56:27 PM	41.862830429
267231	1167752	1893853	2019	07/16/2019 04:17:29 PM	41.864278228
267232	1167746	1893853	2019	06/02/2019 04:09:42 PM	41.864278357
267233	NaN	NaN	2014	06/02/2019 04:09:42 PM	NaN
267234	NaN	NaN	2017	06/02/2019 04:09:42 PM	NaN

	Longitude	Location
267230	-87.717040084	(41.862830429, -87.717040084)
267231	-87.659660218	(41.864278228, -87.659660218)
267232	-87.659682244	(41.864278357, -87.659682244)
267233	NaN	NaN
267234	NaN	NaN

[5 rows x 22 columns]

- To know the progress of work in dask task
- 

#### 0.4 Show the progress of dask work as a bar

```
[5]: from dask.diagnostics import ProgressBar
pbar = ProgressBar()
pbar.register()
```

```
[7]: type(df)
```

```
[7]: dask.dataframe.core.DataFrame
```



```
[6]: %%time
df.count().compute()
```

```
[#####] | 100% Completed | 49.2s
[#####] | 100% Completed | 49.3s
CPU times: user 40.8 s, sys: 4.65 s, total: 45.4 s
Wall time: 49.3 s
```

```
[6]: ID          7054819
Case Number      7054815
Date             7054819
Block            7054819
IUCR             7054819
Primary Type     7054819
Description      7054819
Location Description 7048684
Arrest           7054819
Domestic         7054819
Beat             7054819
District         7054772
Ward             6439989
Community Area   6441324
FBI Code         7054819
X Coordinate     6987226
Y Coordinate     6987226
Year             7054819
Updated On       7054819
Latitude         6987226
Longitude        6987226
Location         6987226
dtype: int64
```

- dask.get: single thread
- dask.threaded.get: multiple thread pool
- dask.multiprocessing.get : multiprocess pool
- 

### 0.5 distributed.Client.get: multiple computer

- 20 CPUs Demo
- 0.6s faster

```
[8]: %%time
df.count().compute(scheduler='processes', num_workers=20) # 20Cpus
```

```
[#####] | 100% Completed | 48.2s
[#####] | 100% Completed | 48.3s
```

CPU times: user 841 ms, sys: 335 ms, total: 1.18 s  
Wall time: 48.6 s

```
[8]: ID 7054819
Case Number 7054815
Date 7054819
Block 7054819
IUCR 7054819
Primary Type 7054819
Description 7054819
Location Description 7048684
Arrest 7054819
Domestic 7054819
Beat 7054819
District 7054772
Ward 6439989
Community Area 6441324
FBI Code 7054819
X Coordinate 6987226
Y Coordinate 6987226
Year 7054819
Updated On 7054819
Latitude 6987226
Longitude 6987226
Location 6987226
dtype: int64
```

- Dask Bags: <https://examples.dask.org/bag.html>
- Dask Bags: <https://docs.dask.org/en/latest/bag.html>
- Dask JSON data: <https://examples.dask.org/applications/json-data-on-the-web.html>
- Dask Tutorial: [https://people.duke.edu/~ccc14/sta-663-2017/18A\\_Dask.html](https://people.duke.edu/~ccc14/sta-663-2017/18A_Dask.html)

### Dask provides 3 data structures

- dask array ~ numpy array
- dask bag ~ Python dictionary
- 

### 0.6 dask dataframe ~ pandas dataframe

```
[1]: import dask
import dask.array as da
import dask.bag as db
import dask.dataframe as dd
from dask import delayed
```

dask arrays

- but break a massive job into tasks that are then executed by a scheduler.
- we can tell the `dask` array how to break the data into chunks for processing.

### **dask data frames**

- treat multiple pandas dataframes that might not simultaneously fit into memory like a single dataframe.

### **dask bags**

- `multiset` is a set that allows repeats. Unlike lists, order is not preserved.

### **Conversion from bag to dataframe**

- `dask_df = dic.to_dataframe(columns=['word', 'n'])`
- `dask_df.head(n=5)`

### **Analyze web-hosted JSON data**

- Ref: <https://examples.dask.org/applications/json-data-on-the-web.html>

### **Read file as a dask bag = db**

- `import dask.bag as db`
- `dic= db.read_text(filenamees).map(json.loads)`
- `dic.take(2)`

### **Dask Gabs**

- `map`, `filter`, `groupby`
- Ref: <https://examples.dask.org/bag.html>

### **Read json as a dask bag**

- `import dask.bag as db`
- `import json`
- `b = db.read_text('data/*.json').map(json.loads)`
- `b.take(2)`

### **Dask bag : Map, Filter, Aggregate**

- `b.filter(lambda record: record['age_variable'] > 30).take(2) # Select only people over 30`
- `b.map(lambda record: record['occupation']).take(2) # Select the occupation field`
- `b.count().compute() # Count total number of records`

## Chain computations

- `result = (b.filter(lambda record: record['age'] > 30) .map(lambda record: record['occupation'])) .frequencies(sort=True) .topk(10, key=1)`
- `result.compute()` # we need to call compute to actually evaluate our result.

## Dask Bag Tutorial 2

- <https://docs.dask.org/en/latest/bag.html>
- `da.stack`, `da.concatenate`, and `da.block`
- Ref: <https://docs.dask.org/en/latest/array-stack.html>

```
[29]: import dask.array as da
```

```
[30]: arr0 = da.from_array(np.zeros((3, 4)), chunks=(1, 2))
      arr1 = da.from_array(np.ones((3, 4)), chunks=(1, 2))
```

```
[31]: arr0
```

```
[31]: dask.array<array, shape=(3, 4), dtype=float64, chunksize=(1, 2),
      chunktype=numpy.ndarray>
```

```
[33]: data = [arr0, arr1]
      data
```

```
[33]: [dask.array<array, shape=(3, 4), dtype=float64, chunksize=(1, 2),
      chunktype=numpy.ndarray>,
      dask.array<array, shape=(3, 4), dtype=float64, chunksize=(1, 2),
      chunktype=numpy.ndarray>]
```

```
[34]: x = da.stack(data, axis=0)
      x.shape
```

```
[34]: (2, 3, 4)
```

```
[35]: x
```

```
[35]: dask.array<stack, shape=(2, 3, 4), dtype=float64, chunksize=(1, 1, 2),
      chunktype=numpy.ndarray>
```

```
[36]: da.stack(data, axis=1).shape
```

```
[36]: (3, 2, 4)
```

## Concatenate

- concatenate existing arrays into a new array, extending them along an existing dimension

```
[38]: import dask.array as da
import numpy as np

arr0 = da.from_array(np.zeros((3, 4)))
arr1 = da.from_array(np.ones((3, 4)))
data = [arr0, arr1]
```

```
[39]: x = da.concatenate(data, axis=0)
x.shape
```

```
[39]: (6, 4)
```

```
[40]: da.concatenate(data, axis=1).shape
```

```
[40]: (3, 8)
```

### Dataframe concante example

- import dask.dataframe as dd
- df = dd.read\_csv(['data/untermaederbrunnen\_station1\_xyz\_intensity\_rgb.txt'], delimiter=',', header=None, names=['x', 'y', 'z', 'intensity', 'r', 'g', 'b'])
- df\_label = dd.read\_csv(['data/untermaederbrunnen\_station1\_xyz\_intensity\_rgb.labels'], header=None, names=['label'])
- df = df.repartition(npartitions=200)
- df\_label = df\_label.repartition(npartitions=200)
- df = dd.concat([df, df\_label], axis=1)

### Dask Groupby

- Ref: <https://examples.dask.org/dataframes/02-groupby.html>

```
[41]: import dask
df = dask.datasets.timeseries()
df
```

```
[41]: Dask DataFrame Structure:
```

	id	name	x	y
npartitions=30				
2000-01-01	int64	object	float64	float64
2000-01-02	...	...	...	...
...	...	...	...	...
2000-01-30	...	...	...	...
2000-01-31	...	...	...	...

Dask Name: make-timeseries, 30 tasks

```
[46]: %%time
df.groupby('name').x.mean().compute(scheduler='processes', num_workers=20)
```

```
[#####] | 100% Completed | 1.2s
[#####] | 100% Completed | 1.3s
CPU times: user 1.05 s, sys: 112 ms, total: 1.17 s
Wall time: 1.34 s
```

```
[46]: name
      Alice      0.001558
      Bob        0.001778
      Charlie    0.001967
      Dan        0.001146
      Edith      0.003717
      Frank      0.001993
      George     -0.000409
      Hannah     0.001473
      Ingrid     0.000094
      Jerry      0.002051
      Kevin      0.000493
      Laura      -0.002057
      Michael    -0.000478
      Norbert    0.000516
      Oliver     -0.001005
      Patricia   0.000314
      Quinn      -0.003217
      Ray        -0.000248
      Sarah      -0.004322
      Tim        0.002823
      Ursula     -0.001512
      Victor     0.002398
      Wendy      0.002392
      Xavier     -0.002486
      Yvonne     0.002242
      Zelda      0.002681
      Name: x, dtype: float64
```

```
[ ]:
```